

---

# FALCON: Few-step Accurate Likelihoods for Continuous Flows

---

Danyal Rehman

Mila - Québec AI Institute

danyal.rehman@mila.quebec

Tara Akhound-Sadegh

Mila - Québec AI Institute

tara.akhoundsadegh@mila.quebec

Artem Gazizov

Harvard Medical School

artem\_gazizov@hms.harvard.edu

Yoshua Bengio

Mila - Québec AI Institute

yoshua.bengio@mila.quebec

Alexander Tong

AITHYRA

atong@aithyra.at

## Abstract

Scalable sampling of molecular states in thermodynamic equilibrium is a long-standing challenge in statistical physics. Boltzmann Generators tackle this problem by pairing a generative model, capable of exact likelihood computation, with importance sampling to obtain consistent samples under the target distribution. Current Boltzmann Generators primarily use continuous normalizing flows (CNFs) trained with flow matching for efficient training of powerful models. However, likelihood calculation for these models is extremely costly, requiring thousands of function evaluations per sample, severely limiting their adoption. In this work, we propose Few-step Accurate Likelihoods for Continuous Flows (FALCON), a method which allows for few-step sampling with a likelihood accurate enough for importance sampling applications by introducing a hybrid training objective that encourages invertibility. We demonstrate that FALCON outperforms state-of-the-art normalizing flow models for molecular Boltzmann sampling and is *two orders of magnitude faster* than the most performant CNFs.

## 1 Introduction

Sampling conformations from the Boltzmann distribution  $p(x) \propto \exp(-\mathcal{E}(x))$  where  $\mathcal{E}(x)$  is the potential energy of a configuration  $x$ , is a foundational challenge in statistical physics. Boltzmann Generators (BGs) have emerged as a way to address this inefficiency by amortizing the cost through the training of a generative model to learn to sample from  $p_\theta(x)$  close to  $p(x)$ . These samples can then be corrected to  $p(x)$  using self-normalized importance sampling (SNIS) [Noé et al., 2019]. SNIS requires efficient access to  $p_\theta(x)$  and  $\mathcal{E}(x)$  for every sample drawn  $x \sim p_\theta(x)$  for the correction step, but guarantees statistical consistency of the corrected samples [Noé et al., 2019, Liu, 2001].

The main design choice in BGs is which type of generator to use. Modern BGs [Klein et al., 2023a, Klein and Noe, 2024] primarily make use of generators based on continuous normalizing flows (CNFs) [Chen et al., 2018, Grathwohl et al., 2019] due to their expressive power, ease of training, and flexibility of parameterization [Köhler et al., 2020] (see Table 1). However, while it is possible to access the  $p_\theta(x)$  of a CNF, it is extremely computationally costly to approximate  $p_\theta(x)$  with sufficient accuracy. Two primary reasons contribute to this cost: (1) Approximate estimators are not sufficiently accurate, making full Jacobian calculations necessary for each step along the flow; and (2) Many steps are necessary to control discretization error for sufficient performance (see Fig. 1).

Recently, there have been significant advancements in few-step generation using flow models [Song et al., 2023, Boffi et al., 2025a, Frans et al., 2025, Sabour et al., 2025, Geng et al., 2025]. These models are powerful few-step generators with flexible architectures and efficient simulation-free training; however, they do not natively admit efficient likelihood estimation, making them unsuitable for the high-precision demands of SNIS, and applications such as Boltzmann Generation [Rehman et al., 2025].

In this work, we investigate how to design a generative model that combines the best of both worlds: the training efficiency and architectural freedom of simulation-free flow models with the fast sampling and likelihood evaluation of discrete-time invertible models. Our main contributions are:

- We introduce Few-step Accurate Likelihoods for Continuous Flows (FALCON): a new continuous flow-based generative model for Boltzmann sampling that is invertible, trained with regression, and supports free-form architectures, while enabling both few-step generation and efficient likelihoods.
- Orthogonally, we introduce a simple and scalable, softly equivariant continuous flow architecture that significantly improves over the current state-of-the-art equivariant flow model architecture.
- We show that FALCON is **two orders of magnitude faster** than CNF-based Boltzmann Generators for equivalent performance (see Fig. 1), drastically reducing the computational cost of CNFs.
- We show that FALCON outperforms the current state-of-the-art normalizing flow-based Boltzmann Generator across all metrics, even when FALCON is given **25× fewer samples** [Tan et al., 2025a].

## 2 Background and Preliminaries

We are interested in drawing statistically independent samples from a target Boltzmann distribution  $p_{\text{target}}$  with partition function  $\mathcal{Z}$ , defined over  $\mathbb{R}^d$ :

$$p_{\text{target}}(x) \propto \exp(-\mathcal{E}(x)), \quad \mathcal{Z} = \int_{\mathbb{R}^d} \exp(-\mathcal{E}(x)) dx \quad (1)$$

where  $\mathcal{E} : \mathbb{R}^d \rightarrow \mathbb{R}$  is the energy of the system, which we can efficiently compute for any  $x$ . In this work, we do not require the energy to be differentiable. Unlike in the pure sampling setting [Akhound-Sadegh et al., 2024, Havens et al., 2025, Akhoud-Sadegh et al., 2025, Midgley et al., 2023, Zhang and Chen, 2022, Vargas et al., 2023], we also assume access to a small biased dataset  $\mathcal{D} = \{x^i\}_{i=1}^N$  of  $N$  samples [Noé et al., 2019]. This makes it possible to perform an initial learning phase that fits a generative model with parameters  $\theta$ , producing a proposal distribution,  $p^\theta(x)$  [Noé et al., 2019].

**Few-step Flow Models.** Flow matching models can require hundreds of steps for accurate approximation of  $p_1$ . To speed up sampling, few-step flow models such as consistency models [Song et al., 2023], optimal transport methods [Pooladian et al., 2023, Tong et al., 2024], and flow map models [Boffi et al., 2025a, Sabour et al., 2025, Geng et al., 2025, Frans et al., 2025, Guo et al., 2025] learn to generate high quality samples in many fewer steps. Recent models that incorporate both current and target sample times have shown flexibility and effectiveness in the few-step regime, where  $u_\theta$  is augmented with the target time  $t$  to capture the average velocity [Geng et al., 2025],  $u(x_s, s, t)(t - s) = \int_s^t v(x_\tau, \tau) d\tau$ , and minimize the following objective:

$$\mathbb{E}_{s,t,x_s} \left[ w(s, t) \left\| u_\theta(x_s, s, t) - \frac{1}{t-s} \int_s^t v(x_\tau, \tau) d\tau \right\|^2 \right], \quad (2)$$

where the average velocity field  $u_\theta$ , parameterized by a neural network, approximates the probability flow ODE  $v$  that transports samples from noise  $p_0$  to data  $p_1$ . At inference, sampling proceeds by discretizing time and updating  $x_{t_i} = x_{t_{i-1}} + (t_i - t_{i-1}) u_\theta(x_{t_{i-1}}, t_{i-1}, t_i)$ . While few-step flow models enable fast generation, they do not provide tractable likelihoods, since  $u_\theta$  need not be invertible unless the training objective is fully optimized (comparable methods are summarized in Table 1).

## 3 Few-step Accurate Likelihoods for Continuous Flows

We now introduce Few-step Accurate Likelihoods for Continuous Flows (FALCON): a novel flow-based generative model designed to address the inherent efficiency limitations of using continuous flow models with Boltzmann Generators. Our method departs from traditional continuous normalizing flow (CNFs) by training a flow map that operates in a few discrete steps, while simultaneously achieving invertibility to ensure fast and accurate likelihood computation for Boltzmann Generation. This is achieved through a hybrid training objective, which, by enabling stable few-step generation, dramatically reduces the inference cost. This efficiency allows us to use much larger and more expressive architectures [Vaswani et al., 2017, Peebles and Xie, 2023, Ma et al., 2024] that were previously computationally infeasible to scale in the Boltzmann Generator setting [Tan et al., 2025a].

Table 1: Related method overview

Method	Invertible	Regression-loss	Few Step	FreeForm Arch.
BioEmu	✗	✓	✗	✓
FlowMaps	✗	✓	✓	✓
TBG	✓	✓	✗	✓
Prose	✓	✗	✓	✗
FFFlows	✓	✗	✓	✓
FALCON (Ours)	✓	✓	✓	✓

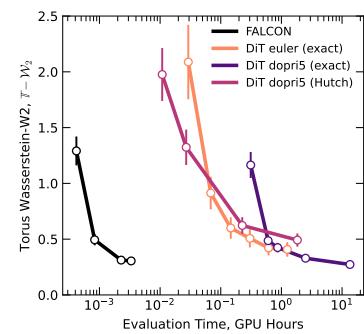


Figure 1: Performance across CNFs during inference using  $10^4$  samples.

**Flow Maps are Flawed Boltzmann Generators.** The core of FALCON is a generative process that learns an invertible map from  $p_0$  to  $p_1$ , in a small number of steps. We first examine the suitability of the existing few-step flows for importance sampling applications, concluding that, on their own, they are not sufficient. We first define the continuous map with respect to a vector field  $v$  as:

$$X_v(x_s, s, t) = \int_s^t v(x_\tau, \tau) d\tau + x_s, \quad (3)$$

and note that under mild regularity conditions on  $v$ , this map is always invertible up to discretization error. For any invertible map, we can compute the change in density with respect to the input using the change of variables formula, at the approximate cost of  $d$  function evaluations, and the determinant, which, while an  $O(d^3)$  operation, is in practice negligible compared to the function evaluation cost.

This invertibility property also holds for flow map models at the optima, which we address in the following proposition.

**Proposition 1.** *Let  $u_\theta^*$  be the minimizer of Eq. 2 with respect to some sufficiently smooth  $v$ . Also, define the discrete flow map:*

$$X_u(x_s, s, t) = x_s + (t - s)u_\theta^*(x_s, s, t) \quad (4)$$

*Then, for any  $(s, t) \in [0, 1]^2$ ,*

1.  $X_u(\cdot, s, t)$  is an invertible map,

2.  $\log p_t^{u^*}(x_t) = \log p_s^{u^*}(x_s) - \log |\det \mathbf{J}_{X_u}(x_s)|$  where  $\mathbf{J}_X = \frac{\partial X}{\partial x_s}$  is the Jacobian of  $X$ .

We provide a more precise statement and proofs for all propositions in Appendix A. This means that *optimal* flow maps are, in some ways, ideal Boltzmann Generators in that they have relatively efficient access to both samples and likelihood; however, this property only holds at the optima, and in the case that  $X_u(\cdot, s, t) = X_v(\cdot, s, t)$  for all  $s, t$ , which in practice is extremely challenging to satisfy.

In practice, for standard flow maps,  $X_u(\cdot, s, t) \neq X_v(\cdot, s, t)$ , meaning  $X_u$  is not guaranteed to be invertible, rendering efficient likelihood computation intractable. However, we note that this condition is actually much stronger than we need. For FALCON, while we would like  $X_u(\cdot, s, t)$  to be close to  $X_v(\cdot, s, t)$ , for efficient likelihood computation, we only require that  $X_u$  is invertible, not that it matches the invertible map defined by  $X_v$ . This leads us to define an additional invertibility loss:

$$\mathcal{L}_{\text{inv}}(\theta) = \mathbb{E}_{s, t, x_s} \|x_s - X_u(X_u(x_s, s, t), t, s)\|^2, \quad (5)$$

to be used in conjunction with the average velocity objective and flow matching objectives,  $\mathcal{L}_{\text{cfm}}$ , for a final loss comprised of three components:

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{cfm}}(\theta) + \lambda_{\text{avg}} \mathcal{L}_{\text{avg}}(\theta) + \lambda_r \mathcal{L}_{\text{inv}}(\theta). \quad (6)$$

Minimizing this loss has a less strict requirement for the correctness of the BG specifically:

**Proposition 2.** *Let  $u_\theta^* \in C^1$  be a minimizer of  $\mathcal{L}_{\text{inv}}$  (Eq. (5)) with respect to some sufficiently regular  $v$ , then for any  $s, t \in [0, 1]^2$ ,  $X_u(\cdot, s, t)$  is an invertible map.*

Thus, minimizing the invertibility loss is sufficient for valid BGs, even without exactly reproducing the continuous-time flow. Note that the proposition provides a constructive guarantee of invertibility, and, in practice, we only require the existence of an inverse, not its explicit form. This condition ensures that FALCON acts as a consistent generator of the target energy distribution,  $\mathcal{E}(x)$ , while benefiting from fast inference-time scalability. We detail the FALCON training algorithm in Appendix F.

**FALCON Enables Scalable Architectures.** Previous Boltzmann Generators based on CNFs adopted small equivariant architectures [Klein et al., 2023a, Klein and Noe, 2024, Tan et al., 2025a, Aggarwal et al., 2025] up to 2.3 million parameters. These models are limited in their scale due to the high cost of inference with multi-step adaptive step size samplers, which are needed to control the error in the likelihood. FALCON, by enabling relatively cheap few-step sampling, allows us to use larger networks, which we test up to 88.7 million parameters, significantly improving performance.

**Efficient Implementation.** As noted in previous works,  $\mathcal{L}_1$  can be efficiently implemented using a single Jacobian vector product (JVP) call using forward automatic differentiation [Geng et al., 2025]:

$$u_\theta(x_s, s, t), \frac{du_\theta}{ds} = \text{jvp}(u_\theta, (x_s, s, t), (v_s, 1, 0))$$

where the `jvp` function takes a callable function, inputs, and a vector which is the vector part of the JVP. Additionally, for this loss specifically, we can combine  $\mathcal{L}$  and  $\mathcal{L}_{\text{avg}}$  if we implement  $v(x_s, s) = u_\theta(x_s, s, s)$ , i.e. passing the same time to  $u$ , representing the instantaneous velocity.

Our method is the first to require flow maps both in the forward and backward directions. We therefore must consider  $u_\theta(x_s, t, s)$ , i.e. the backwards parameterization, specifically at the discontinuity when  $s = t$ . When  $t \rightarrow s^+$ , then  $u_\theta(x_s, s, t) = v(x_s, s)$ , but when  $t \rightarrow s^-$ , then  $u_\theta(x_s, s, t) = -v(x_s, s)$ . We address this by parameterizing our flow map  $u_\theta$  such that  $u_\theta(x_s, s, t) = \text{sign}(t - s)h_\theta(x_s, s, t)$ .

## 4 Experiments

In this section, we first demonstrate that FALCON achieves more scalable performance over state-of-the-art CNFs across both global and local metrics on tri-alanine, alanine tetrapeptide, and hexa-alanine (Table 3). Next, we show that FALCON exceeds the performance of state-of-the-art discrete NFs, even when they are given vastly larger sampling budgets (Fig. 3). Lastly, we elucidate the role of regularization in both achieving invertibility and enhancing generative performance (Fig. 4).

### 4.1 FALCON Outperforms State-of-the-Art Methods

**Superior Scalability Over Continuous Flows.** Computing likelihoods in CNFs is computationally prohibitive, limiting their scalability in the BG setting. Although the current state-of-the-art, ECNF++, performs exceptionally well on ESS and  $\mathbb{T}\mathcal{W}_2$  for alanine dipeptide (see Table 2) [Tan et al., 2025a], it fails to scale to larger molecules, as seen in Table 3. In contrast, for larger systems—tri-alanine, alanine tetrapeptide, and hexa-alanine—FALCON substantially outperforms ECNF++ across all metrics, demonstrating superior scalability to larger molecular systems. The true MD energy distributions, learned proposals, and re-sampled energies for alanine dipeptide, tri-alanine, alanine tetrapeptide, and hexa-alanine are all shown in Fig. 2.

Table 3: Quantitative results on tri-alanine (AL3), alanine tetrapeptide (AL4), and hexa-alanine (AL6). Baseline methods presented with SNIS, unless stated otherwise. Evaluations are conducted over  $10^4$  samples.

Algorithm ↓	Tri-alanine (AL3)			Tetrapeptide (AL4)			Hexa-alanine (AL6)		
	ESS ↑	$\mathcal{E}\mathcal{W}_2 \downarrow$	$\mathbb{T}\mathcal{W}_2 \downarrow$	ESS ↑	$\mathcal{E}\mathcal{W}_2 \downarrow$	$\mathbb{T}\mathcal{W}_2 \downarrow$	ESS ↑	$\mathcal{E}\mathcal{W}_2 \downarrow$	$\mathbb{T}\mathcal{W}_2 \downarrow$
ECNF++	0.003 ± 0.002	2.206 ± 0.813	0.962 ± 0.253	0.016 ± 0.001	5.638 ± 0.483	1.002 ± 0.061	0.006 ± 0.001	10.668 ± 0.285	1.902 ± 0.055
SBG IS	0.052 ± 0.013	0.758 ± 0.506	0.502 ± 0.16	0.046 ± 0.014	1.068 ± 0.495	0.969 ± 0.067	0.034 ± 0.015	<b>1.021 ± 0.239</b>	1.431 ± 0.085
SBG SMC	—	0.598 ± 0.084	0.503 ± 0.029	—	1.007 ± 0.382	1.039 ± 0.069	—	1.189 ± 0.357	1.444 ± 0.140
FALCON (Ours)	<b>0.074 ± 0.017</b>	<b>0.557 ± 0.136</b>	<b>0.450 ± 0.035</b>	<b>0.052 ± 0.016</b>	<b>0.625 ± 0.102</b>	<b>0.851 ± 0.088</b>	<b>0.065 ± 0.007</b>	1.903 ± 0.122	<b>1.255 ± 0.105</b>

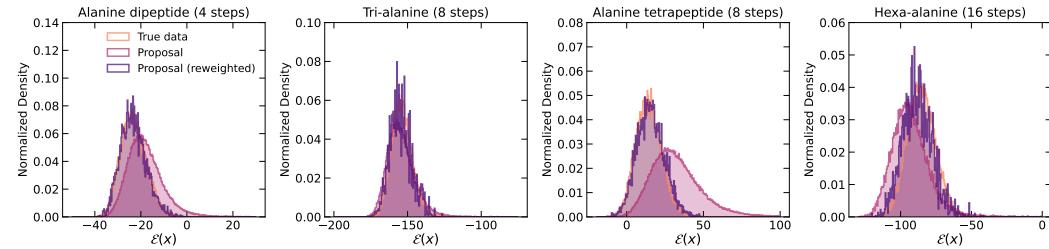


Figure 2: True MD energy distribution with best FALCON unweighted and re-sampled proposals for alanine dipeptide (left), tri-alanine (center left), and alanine tetrapeptide (center right), and hexa-alanine (right).

**Improved Sample Quality Over Discrete Flows.** Discrete normalizing flows have recently been shown to be the most performant Boltzmann Generators [Tan et al., 2025a,b]. SBG [Tan et al., 2025a], based on the TARFlow architecture [Zhai et al., 2024], outperforms all previously reported methods across both global and local metrics (see Table 3). Here, we demonstrate that FALCON, even in a few steps, can outperform SBG across all metrics. We make two main claims to assert FALCON’s competitive advantage, in comparison with discrete NFs: (1) Discrete NFs, despite being fast one-step generators, consistently underperform compared to FALCON across all global and local metrics (Table 3); and (2) Increasing the number of samples can partially close this gap; however, even with  $5 \times 10^6$  samples— $250\times$  more than those used to evaluate FALCON—SBG’s performance on  $\mathcal{E}\mathcal{W}_2$  remains significantly worse than that of a 4-step FALCON Flow (as demonstrated in Fig. 3).

## 4.2 Analysis of Computational Efficiency

**Training Efficiency Compared to Discrete NFs.** Discrete NFs benefit from fast inference, but are slow and unstable to train due to the maximum likelihood objective [Xu and Campbell, 2023, Andrade, 2024]. By contrast, CNFs trained with a flow matching objective trade more stable and faster training for slower inference. When considering both training and evaluation time together (Table 4), we see that FALCON—despite being marginally slower at inference than the discrete NFs for the same number of samples—achieves faster cumulative training + inference times for superior performance due to the expedited training objective.

**Inference Efficiency Compared to CNFs.** A primary contribution of FALCON is the dramatic reduction in the computational cost required to achieve high-quality samples with accurate likelihoods. Fig. 1 directly illustrates this advantage. To reach a comparable level of performance on the  $\mathbb{T}\text{-}\mathcal{W}_2$  metric for alanine dipeptide, a traditional CNF requires inference times that are two orders of magnitude longer than FALCON. This efficiency gain is what enables the use of larger, more expressive architectures and makes large-scale molecular sampling practical. For additional discussion, see Appendix G.3.

Table 4: Cumulative training + inference time across flows.  $10^4$  samples evaluated with ( $\text{atol} = \text{rtol} = 10^{-5}$  for our CNF and 4-step FALCON). All experiments were conducted on one NVIDIA L40S with batch size 1024. Note: For hexa-alanine, obtaining  $10^4$  samples from the Dopri5-integrated CNF was computationally infeasible.

	ECNF++	SBG	DiT CNF (Ours)	FALCON (Ours)
Alanine dipeptide	12.52	16.83	11.21	7.65
Tri-alanine	19.59	24.67	20.67	18.31
Alanine tetrapeptide	32.17	41.67	41.44	38.14
Hexa-alanine	137.4	57.5	—	60.4

## 4.3 Ablation Studies

### Verifying FALCON’s Invertibility.

As CNFs are only invertible at convergence, we introduce a regularization term in the loss to promote numerical invertibility in the few-step regime. In Fig. 4, we demonstrate the trade-off from this term on both the ESS and 2-Wasserstein distance on dihedral angles: weak regularization leads to poor invertibility and degraded performance, whereas strong regularization enforces flow invertibility, albeit at the cost of reduced sample quality. We fix the regularization constant to 10.0 for all experiments performed, unless stated otherwise to balance performance and numerical invertibility.

## 5 Conclusion

In this work, we present Few-step Accurate Likelihoods for Continuous Flows, a novel few-step flow-based generative model that enables scalable and efficient Boltzmann distribution sampling by combining the expressiveness and training efficiency of modern flow-based architectures with rapid sampling and accurate likelihood estimation. Through a hybrid training objective, FALCON overcomes the computational bottlenecks of likelihood evaluation that have hindered the adoption of Boltzmann Generators. Empirically, FALCON outperforms state-of-the-art discrete normalizing flows while achieving two orders of magnitude faster performance than comparably accurate CNF-based Boltzmann Generators, thereby making large-scale molecular sampling tasks more practical and advancing the applicability of Boltzmann Generators in domains such as drug discovery and materials science.

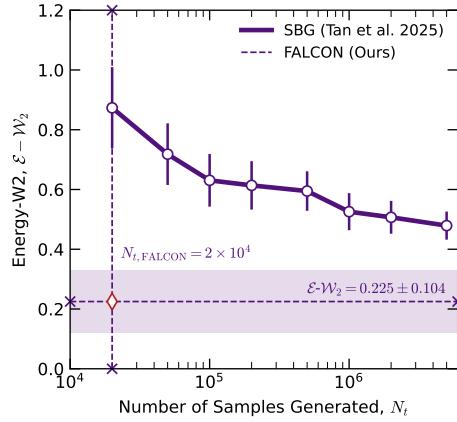


Figure 3: Performance with additional samples.

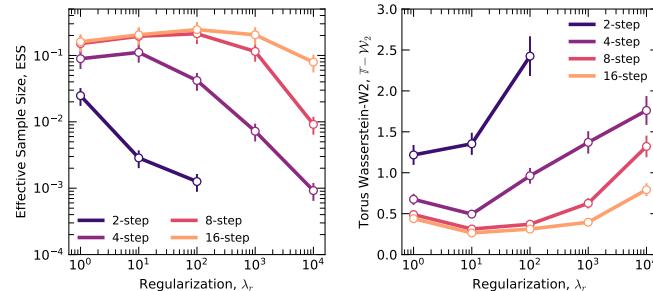


Figure 4: Performance trade-off with increasing regularization.

## References

- J. Abramson, J. Adler, J. Dunger, R. Evans, T. Green, A. Pritzel, O. Ronneberger, L. Willmore, A. J. Ballard, J. Bambrick, S. W. Bodenstein, D. A. Evans, C.-C. Hung, M. O'Neill, D. Reiman, K. Tunyasuvunakool, Z. Wu, A. Žemgulytė, E. Arvaniti, C. Beattie, O. Bertolli, A. Bridgland, A. Cherepanov, M. Congreve, A. I. Cowen-Rivers, A. Cowie, M. Figurnov, F. B. Fuchs, H. Gladman, R. Jain, Y. A. Khan, C. M. R. Low, K. Perlin, A. Potapenko, P. Savy, S. Singh, A. Stecula, A. Thillaisundaram, C. Tong, S. Yakneen, E. D. Zhong, M. Zielinski, A. Žídek, V. Bapst, P. Kohli, M. Jaderberg, D. Hassabis, and J. M. Jumper. Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*, 630(8016):493–500, June 2024. ISSN 1476-4687. doi: 10.1038/s41586-024-07487-w. URL <https://www.nature.com/articles/s41586-024-07487-w>. Publisher: Nature Publishing Group.
- R. Aggarwal, J. Chen, N. M. Boffi, and D. R. Koes. BoltzNCE: Learning likelihoods for boltzmann generation with stochastic interpolants and noise contrastive estimation. *arXiv preprint arXiv:2507.00846*, 2025.
- T. Akhound-Sadegh, J. Rector-Brooks, A. J. Bose, S. Mittal, P. Lemos, C.-H. Liu, M. Sendera, S. Ravanbakhsh, G. Gidel, Y. Bengio, N. Malkin, and A. Tong. Iterated Denoising Energy Matching for Sampling from Boltzmann Densities, Feb. 2024. URL <http://arxiv.org/abs/2402.06121>. arXiv:2402.06121 [cs, stat].
- T. Akhound-Sadegh, J. Lee, A. J. Bose, V. D. Bortoli, A. Doucet, M. M. Bronstein, D. Beaini, S. Ravanbakhsh, K. Neklyudov, and A. Tong. Progressive inference-time annealing of diffusion models for sampling from boltzmann densities. In *Advances in Neural Information Processing Systems*, 2025.
- D. Andrade. Stable training of normalizing flows for high-dimensional variational inference. *arXiv preprint arXiv:2402.16408*, 2024.
- N. M. Boffi, M. S. Albergo, and E. Vanden-Eijnden. How to build a consistency model: Learning flow maps via self-distillation, 2025a. URL <https://arxiv.org/abs/2505.18825>.
- N. M. Boffi, M. S. Albergo, and E. Vanden-Eijnden. Flow map matching with stochastic interpolants: A mathematical framework for consistency models. *Transactions on Machine Learning Research*, 2025b.
- R. T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. K. Duvenaud. Neural ordinary differential equations. *Neural Information Processing Systems (NIPS)*, 2018.
- J. R. Dormand and P. J. Prince. Runge-kutta triples. *Computers & Mathematics with Applications*, 12(9):1007–1017, 1986.
- F. Draxler, P. Sorrenson, L. Zimmermann, A. Rousselot, and U. Köthe. Free-form flows: Make any architecture a normalizing flow. In *International Conference on Artificial Intelligence and Statistics*, pages 2197–2205. PMLR, 2024.
- K. Frans, D. Hafner, S. Levine, and P. Abbeel. One step diffusion via shortcut models, 2025. URL <https://arxiv.org/abs/2410.12557>.
- Z. Geng, A. Pokle, W. Luo, J. Lin, and J. Z. Kolter. Consistency models made easy. *arXiv preprint arXiv:2406.14548*, 2024.
- Z. Geng, M. Deng, X. Bai, J. Z. Kolter, and K. He. Mean flows for one-step generative modeling, 2025. URL <https://arxiv.org/abs/2505.13447>.
- W. Grathwohl, R. T. Q. Chen, J. Bettencourt, I. Sutskever, and D. Duvenaud. FFJORD: free-form continuous dynamics for scalable reversible generative models. In *International Conference on Representation Learning (ICLR)*, 2019.
- Y. Guo, W. Wang, Z. Yuan, R. Cao, K. Chen, Z. Chen, Y. Huo, Y. Zhang, Y. Wang, S. Liu, and Y. Wang. Splitmeanflow: Interval splitting consistency in few-step generative modeling, 2025. URL <https://arxiv.org/abs/2507.16884>.

- A. Havens, B. K. Miller, B. Yan, C. Domingo-Enrich, A. Sriram, B. Wood, D. Levine, B. Hu, B. Amos, B. Karrer, X. Fu, G.-H. Liu, and R. T. Q. Chen. Adjoint sampling: Highly scalable diffusion samplers via adjoint matching, 2025. URL <https://arxiv.org/abs/2504.11713>.
- B. Jing, S. Eismann, P. Suriana, R. J. Townshend, and R. Dror. Learning from protein structure with geometric vector perceptrons. *arXiv preprint arXiv:2009.01411*, 2020.
- T. Karras, M. Aittala, T. Aila, and S. Laine. Elucidating the design space of diffusion-based generative models. *Advances in neural information processing systems*, 35:26565–26577, 2022.
- L. Kish. Confidence intervals for clustered samples. *American Sociological Review*, 22(2):154–165, 1957.
- L. Klein and F. Noe. Transferable Boltzmann Generators. *Advances in Neural Information Processing Systems*, 37:45281–45314, Dec. 2024. URL [https://proceedings.neurips.cc/paper\\_files/paper/2024/hash/5035a409f5798e188079e236f437e522-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/5035a409f5798e188079e236f437e522-Abstract-Conference.html).
- L. Klein, A. Krämer, and F. Noé. Equivariant flow matching. *Advances in Neural Information Processing Systems*, 36:59886–59910, Dec. 2023a. URL [https://proceedings.neurips.cc/paper\\_files/paper/2023/hash/bc827452450356f9f558f4e4568d553b-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/bc827452450356f9f558f4e4568d553b-Abstract-Conference.html).
- L. Klein, A. Krämer, and F. Noé. Equivariant flow matching. *Neural Information Processing Systems (NeurIPS)*, 2023b.
- J. Köhler, L. Klein, and F. Noé. Equivariant flows: exact likelihood generative learning for symmetric densities. *International Conference on Machine Learning (ICML)*, 2020.
- J. Köhler, M. Invernizzi, P. D. Haan, and F. Noe. Rigid Body Flows for Sampling Molecular Crystal Structures. In *Proceedings of the 40th International Conference on Machine Learning*, pages 17301–17326. PMLR, July 2023. URL <https://proceedings.mlr.press/v202/kohler23a.html>. ISSN: 2640-3498.
- Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le. Flow matching for generative modeling. *International Conference on Learning Representations (ICLR)*, 2023.
- J. S. Liu. *Monte Carlo Strategies in Scientific Computing*. Springer, 2001.
- N. Ma, M. Goldstein, M. S. Albergo, N. M. Boffi, E. Vanden-Eijnden, and S. Xie. Sit: Exploring flow and diffusion-based generative models with scalable interpolant transformers. In *Computer Vision – ECCV 2024*, 2024.
- L. I. Midgley, V. Stimper, G. N. Simm, B. Schölkopf, and J. M. Hernández-Lobato. Flow annealed importance sampling bootstrap. *International Conference on Learning Representations (ICLR)*, 2023.
- F. Noé, S. Olsson, J. Köhler, and H. Wu. Boltzmann generators: Sampling equilibrium states of many-body systems with deep learning. *Science*, 365(6457):eaaw1147, 2019.
- W. Peebles and S. Xie. Scalable Diffusion Models with Transformers. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 4172–4182, Paris, France, Oct. 2023. IEEE. ISBN 9798350307184. doi: 10.1109/ICCV51070.2023.00387. URL <https://ieeexplore.ieee.org/document/10377858/>.
- A.-A. Pooladian, H. Ben-Hamu, C. Domingo-Enrich, B. Amos, Y. Lipman, and R. T. Q. Chen. Multisample flow matching: Straightening flows with minibatch couplings. In *International Conference on Machine Learning*, 2023.
- D. Rehman, O. Davis, J. Lu, J. Tang, M. Bronstein, Y. Bengio, A. Tong, and A. J. Bose. Fort: Forward-only regression training of normalizing flows, 2025. URL <https://arxiv.org/abs/2506.01158>.
- A. Rizzi, P. Carloni, and M. Parrinello. Free energies at qm accuracy from force fields via multimap targeted estimation. *Proceedings of the National Academy of Sciences*, 2023.

- A. Sabour, S. Fidler, and K. Kreis. Align your flow: Scaling continuous-time flow map distillation, 2025. URL <https://arxiv.org/abs/2506.14603>.
- M. Schebek, M. Invernizzi, F. Noé, and J. Rogal. Efficient mapping of phase diagrams with conditional boltzmann generators. *Machine Learning: Science and Technology*, 2024.
- Y. Song and P. Dhariwal. Improved techniques for training consistency models. *arXiv preprint arXiv:2310.14189*, 2023.
- Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon, and B. Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations (ICLR)*, 2021a.
- Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon, and B. Poole. Score-based generative modeling through stochastic differential equations, 2021b. URL <https://arxiv.org/abs/2011.13456>.
- Y. Song, P. Dhariwal, M. Chen, and I. Sutskever. Consistency models. In *International Conference on Machine Learning*, 2023.
- P. Sorrenson, F. Draxler, A. Rousselot, S. Hummerich, L. Zimmermann, and U. Köthe. Lifting architectural constraints of injective flows. In *International Conference on Learning Representations*, 2024.
- C. B. Tan, A. J. Bose, C. Lin, L. Klein, M. M. Bronstein, and A. Tong. Scalable Equilibrium Sampling with Sequential Boltzmann Generators, Feb. 2025a. URL <http://arxiv.org/abs/2502.18462> [cs]. arXiv:2502.18462 [cs].
- C. B. Tan, M. Hassan, L. Klein, S. Syed, D. Beaini, M. M. Bronstein, A. Tong, and K. Neklyudov. Amortized sampling with transferable normalizing flows. In *Advances in Neural Information Processing Systems*, 2025b.
- A. Tong, K. Fatras, N. Malkin, G. Huguet, Y. Zhang, J. Rector-Brooks, G. Wolf, and Y. Bengio. Improving and generalizing flow-based generative models with minibatch optimal transport. *Transactions on Machine Learning Research*, 2024.
- F. Vargas, W. Grathwohl, and A. Doucet. Denoising diffusion samplers. *International Conference on Learning Representations (ICLR)*, 2023.
- A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need, 2017.
- P. Wirnsberger, A. J. Ballard, G. Papamakarios, S. Abercrombie, S. Racanière, A. Pritzel, D. Jimenez Rezende, and C. Blundell. Targeted free energy estimation via learned mappings. *J. Chem. Phys.*, 2020.
- Z. Xu and T. Campbell. Embracing the chaos: analysis and diagnosis of numerical instability in variational flows. *Advances in Neural Information Processing Systems*, 36:32360–32386, 2023.
- S. Zhai, R. Zhang, P. Nakkiran, D. Berthelot, J. Gu, H. Zheng, T. Chen, M. A. Bautista, N. Jaitly, and J. Susskind. Normalizing Flows are Capable Generative Models, Dec. 2024. URL <http://arxiv.org/abs/2412.06329>. arXiv:2412.06329 [cs].
- Q. Zhang and Y. Chen. Path integral sampler: a stochastic control approach for sampling. *International Conference on Learning Representations (ICLR)*, 2022.

## Appendix

### A Theory

**Proposition 1.** Let  $u_\theta^*$  be the minimizer of Eq. 2 with respect to some sufficiently smooth  $v$ . Also, define the discrete flow map:

$$X_u(x_s, s, t) = x_s + (t - s)u_\theta^*(x_s, s, t) \quad (4)$$

Then, for any  $(s, t) \in [0, 1]^2$ ,

1.  $X_u(\cdot, s, t)$  is an invertible map,
2.  $\log p_t^{u^*}(x_t) = \log p_s^{u^*}(x_s) - \log |\det \mathbf{J}_{X_u}(x_s)|$  where  $\mathbf{J}_X = \frac{\partial X}{\partial x_s}$  is the Jacobian of  $X$ .

We first restate the proposition with what we mean by sufficiently smooth. In particular, we require  $v$  to be uniformly Lipschitz continuous in  $x$  and continuous in  $t$ . We note that if the Lipschitz condition does not hold then the Peano existence theorem implies that the initial value problem (IVP) of the ODE may not be invertible. If, however, we have both continuity in  $x$  and  $t$  as well as a Lipschitz condition for  $v$  on  $x$ , then by the Picard-Lindelöf Theorem, we have existence and uniqueness of the IVP. We first recall Picard-Lindelöf.

**Theorem 1** (Picard-Lindelöf). *Let  $D \subset \mathbb{R}^d \times \mathbb{R}$  be a closed rectangle with  $(x_0, t_0) \in \text{int}D$ , the interior of  $D$ . Let  $v : D \rightarrow \mathbb{R}^d$  be a function that is continuous in  $t$  and Lipschitz continuous in  $x$  (with Lipschitz constant independent from  $t$ ). Then there exists some  $\epsilon > 0$  such that the initial value problem is:*

$$\frac{dx}{dt} = v(x_t, t), \quad x(t_0) = x_0 \quad (7)$$

has a unique solution  $x(t)$  on the interval  $[t_0 - \epsilon, t_0 + \epsilon]$ .

Under these conditions we are now able to prove the proposition.

*Proof.* Recall that if  $u_\theta^*$  minimizes Eq. 2:

$$\mathbb{E}_{s,t,x_s} \left[ w(s, t) \|u_\theta(x_s, s, t) - \frac{1}{t-s} \int_s^t v(x_\tau, \tau) d\tau\|^2 \right]$$

then we have that  $u_\theta^*(x_s, s, t) = \frac{1}{t-s} \int_s^t v(x_\tau, \tau) d\tau$  for all  $s, t \in [0, 1]^2$ . Furthermore we have

$$X_u(x_s, s, t) = x_s + (t - s)u_\theta^*(x_s, s, t) \quad (8)$$

$$= x_s + (t - s) \frac{1}{t - s} \int_s^t v(x_\tau, \tau) d\tau \quad (9)$$

$$= x_s + \int_s^t v(x_\tau, \tau) d\tau \quad (10)$$

which by application of the Picard-Lindelöf theorem is invertible for all  $s, t$ , proving part 1 of the proposition.

To prove part 2, we note that Eq. 10 is differentiable with respect to  $x_s$  and  $X_s$  is invertible, therefore by the change-of-variables formula we arrive at part 2 of the proposition.  $\square$

**Proposition 2.** Let  $u_\theta^* \in C^1$  be a minimizer of  $\mathcal{L}_{\text{inv}}$  (Eq. (5)) with respect to some sufficiently regular  $v$ , then for any  $s, t \in [0, 1]^2$ ,  $X_u(\cdot, s, t)$  is an invertible map.

*Proof.* First we recall Eq. 5:

$$\mathcal{L}_{\text{inv}} = \mathbb{E}_{s,t,x_s} \|x_s - X_u(X_u(x_s, s, t), t, s)\|^2 \quad (11)$$

In this case, the minimizer of  $\mathcal{L}_{\text{inv}}$  that is also in  $C^1$  is pointwise invertible by definition as  $X_u(\cdot, t, s)$  is the inverse of  $X_u(\cdot, s, t)$ .  $\square$

We note that this loss ensures *pointwise* invertibility. Additional restrictions are needed to ensure that the log likelihood can be calculated efficiently.

## B Other formulations for $\mathcal{L}_{\text{avg}}$

There are various formulations of  $\mathcal{L}_{\text{avg}}$  that have been explored in the recent literature. We note that FALCON directly benefits from any new advancements in this rapidly-evolving research direction. Specifically, the following three different  $\mathcal{L}_{\text{avg}}$  losses can be used for training flow maps.

1.  $\mathcal{L}_1 \triangleq \mathbb{E}_{s,t,x_s} \left\| u_\theta(s, t, x_s) - \text{sg}\left(v(x_s, s) - (t-s)(v(x_s, s)\partial_{x_s} u_\theta + \partial_s u_\theta)\right) \right\|^2$ , which is equivalent to the MeanFlow loss of [Geng et al. \[2025\]](#), as well as the ESD objective in [Boffi et al. \[2025a\]](#). Note that since  $x_s = sx_1 + (1-s)x_0$ , we can directly use  $v(x_s, s) = x_1 - x_0$ .
2.  $\mathcal{L}_2 \triangleq \mathbb{E}_{s,t,x_s} \left\| u_\theta(s, t, x_s) - \text{sg}\left(\lambda u_\theta(x_s, s, r) + (1-\lambda)u_\theta(X_u(x_s, s, r), r, t)\right) \right\|^2$ , for  $r = \lambda s + (1-\lambda)t$ , is equivalent to the SplitMeanFlow loss of [Guo et al. \[2025\]](#), as well as the scaled PSD objective in [Boffi et al. \[2025a\]](#). In [Boffi et al. \[2025a\]](#), two ways of sampling the intermediate time  $r$  are explored. Deterministic mid-point sampling, which sets  $\lambda = 0.5$ , as well as uniform sampling of  $\lambda$ . [Guo et al. \[2025\]](#) only explores uniform sampling. This is also a generalization of [Frans et al. \[2025\]](#), which additionally samples  $s$  and  $t$  in a binary-tree fashion.
3.  $\mathcal{L}_3 \triangleq \mathbb{E}_{s,t,x_s} \left\| \partial_t X_u(x_s, s, t) - \text{sg}\left(u_\theta(X_u(x_s, s, t), t, t)\right) \right\|^2$  is the Lagrangian objective presented in [Boffi et al. \[2025a\]](#).

As discussed in Section 3, we choose  $\mathcal{L}_{\text{avg}} \triangleq \mathcal{L}_1$  in this paper. This choice is based on the superior performance of this loss compared to  $\mathcal{L}_3$  in the image experiments of [Sabour et al. \[2025\]](#) and the fact that using jvp, it can be trained more efficiently compared to  $\mathcal{L}_2$ .

## C Extended Related Work

**Boltzmann Generators.** Boltzmann Generators (BGs) [[Noé et al., 2019](#)] are used to sample molecular conformations [[Klein et al., 2023a](#)] and enable consistent estimates of thermodynamic observables [[Wirsberger et al., 2020](#), [Rizzi et al., 2023](#), [Schebek et al., 2024](#)]. While traditionally BGs are based on discrete normalizing flows, more recent work in machine learning makes use of more powerful continuous normalizing flow architectures for invariance [[Köhler et al., 2020](#), [Köhler et al., 2023](#)] and expressive power [[Klein et al., 2023a](#), [Klein and Noe, 2024](#)]. A few other works have explored the usage of approximate likelihoods [[Draxler et al., 2024](#), [Sorrenson et al., 2024](#), [Aggarwal et al., 2025](#)], but have until now been unable to scale.

**Few-step Flows.** Diffusion and flow matching methods are now central to domains from vision [[Song et al., 2021a,b](#), [Lipman et al., 2023](#)] to scientific applications in material and drug discovery [[Abramson et al., 2024](#), [Noé et al., 2019](#)]. Scalable regression-based losses make these models fast to train, yet inference remains costly due to the need for numerous vector field integrations, motivating efforts to reduce computational expense. One- and few-step methods, like consistency models [[Song et al., 2023](#), [Song and Dhariwal, 2023](#), [Geng et al., 2024](#)], shortcut models [[Frans et al., 2025](#)], and flow maps [[Boffi et al., 2025a,b](#), [Sabour et al., 2025](#), [Guo et al., 2025](#)] have gained recent attention. Although there have been numerous efforts in improving generative performance in image settings, developing invertible few-step flows for scientific applications has seen far less interest. Our work, as far as we know, is the first demonstration of an invertible few-step flow with fast likelihoods.

**Relation to Free-Form Flows** Free-form Flows (FFF) enable arbitrary architectures to function as normalizing flows by jointly training an auxiliary network to approximate the inverse of the forward model [[Draxler et al., 2024](#)]. The forward model maps data to latents, while the auxiliary network regularizes this mapping by learning a tractable estimator for the Jacobian term in the change of variables. As correctly highlighted by [Draxler et al. \[2024\]](#), although the loss encourages invertibility, this guarantee only holds when the reconstruction error is sufficiently small—a condition that can be difficult to meet in practice. In our setting, FALCON is trained with an entirely different loss function—albeit with a cycle-consistency term to promote invertibility—similar to that observed by [Draxler et al. \[2024\]](#). In addition, we explicitly validate invertibility for FALCON, in line with their observations through a dedicated experiment.

## D Experimental Details

All training experiments are run on a heterogeneous cluster of NVIDIA H100 and L40S GPUs using distributed data parallelism (DDP). All models were trained with three random seeds, and reported values are averages across runs. Additional training and inference details are included in the following. For benchmarks with existing methods, in cases where dashes are present, data was unavailable, except for the case with SBG SMC [Tan et al., 2025a], where ESS is not a valid metric.

### D.1 Training Details

**Architecture.** We adopt a Diffusion Transformer (DiT) backbone for FALCON, with model size scaled to the complexity of the molecular system under consideration. The details of the backbone configuration are in Table 5 below.

Table 5: Overview of FALCON configurations across datasets.

Dataset	Hidden Size	Blocks	Heads	Cond. Dim.	Parameters (M)
Alanine dipeptide	192	6	6	64	3.2
Tri-alanine	384	12	12	64	23.2
Alanine tetrapeptide	384	12	12	64	23.2
Hexa-alanine	768	12	12	64	88.7

**Training Configuration** All models were trained with an exponential moving average (EMA) on the weights using a decay rate of 0.999. Logit values were clipped at 0.002, with compositional energy regularization disabled. For evaluation, we generated  $10^4$  proposal samples, and used the same number for re-sampling and computing all metrics. Lastly, center-of-mass augmentation was applied with a standard deviation of  $1/\sqrt{n}$ , where  $n$  is the number of particles.

**Hyperparameters** Optimization was performed using AdamW with learning rate  $\text{lr} = 5 \times 10^{-4}$ ,  $\beta = (0.9, 0.999)$ ,  $\epsilon = 10^{-8}$ , and weight decay  $10^{-4}$ . A cosine annealing learning rate schedule with a warm-up phase covering 5% of the training iterations was also used.

**Datasets.** We evaluate the performance of FALCON on equilibrium conformation sampling tasks, focusing on alanine dipeptide (ALDP), tri-alanine (AL3), alanine tetrapeptide (AL4), and hexa-alanine (AL6). Datasets are obtained from implicit solvent molecular dynamics (MD) simulations with the amber-14 force field, as detailed in Appendix D.3. We train on biased data and test on a held-out unbiased dataset, using self-normalized importance sampling (SNIS) and force-field energies to compute log-likelihoods and re-sample molecules from the target Boltzmann density.

**Baselines.** We benchmark FALCON against both discrete and continuous normalizing flows. We include three discrete normalizing flow baselines: (1) SE(3)-EACF [Midgley et al., 2023]; (2) SBG [Tan et al., 2025a] with standard SNIS (SBG IS); and (3) SBG with SMC sampling (SBG SMC), as well as three continuous flows: (1) ECNF [Klein et al., 2023b]; (2) ECNF++ [Tan et al., 2025a]; and (3) BoltzNCE [Aggarwal et al., 2025], a recent SE(3)-equivariant architecture leveraging geometric vector perceptrons (GVPs) [Jing et al., 2020] on alanine dipeptide. For all continuous flows, samples and likelihoods are generated by integrating over the vector field using the Dormand–Prince 4(5) integrator with  $\text{atol} = \text{rtol} = 10^{-5}$  [Dormand and Prince, 1986] to ensure a fair comparison between methods. More details on architectures and parameters are covered in Appendix D.1.

**Metrics.** We report Effective Sample Size (ESS), and the 2-Wasserstein distance on both the energy distribution ( $\mathcal{E}-\mathcal{W}_2$ ), and dihedral angles ( $\mathbb{T}-\mathcal{W}_2$ ). Full definitions are included in Appendix E. The energy captures local details, as minor atomic displacements yield large variations in the energy distribution, while  $\mathbb{T}-\mathcal{W}_2$  captures global structure via mode coverage across metastable states. We include energy histograms in the main text, with Ramachandran plots relegated to Appendix G.4. For robustness, all quantitative experiments are performed on three seeds of the model and reported as mean  $\pm$  standard deviation in the tables and figures. For all benchmarks, in cases where dashes are present, data was unavailable, except for SBG SMC [Tan et al., 2025a], where ESS is not a valid metric.

### D.2 Inference Schedulers

We evaluated five inference schedules to assess their impact on generative performance. The linear baseline distributes steps uniformly across the trajectory, while the geometric, cosine, Chebyshev,

and EDM schedules bias step allocation to emphasize regions where the data distribution is more sensitive. The mathematical definitions and parameter settings for each schedule are provided below.

**Linear.** The linear schedule spaces is distributed uniformly between 1 and 0, with  $N \in \mathbb{N}$  denoting the total number of inference steps such that:

$$t_i = 1 - \frac{i}{N}, \quad i \in \{0, \dots, N\}. \quad (12)$$

**Geometric.** The geometric schedule exponentially allocates more resolution to early steps. We let  $\alpha \in \mathbb{R}$  and  $\alpha > 1$ , denote the geometric base. For all experiments conducted, we set  $\alpha = 2$ :

$$t_i = \frac{\alpha^{N-i} - 1}{\alpha^N - 1}, \quad i \in \{0, \dots, N\}. \quad (13)$$

**Cosine.** The cosine schedule follows a squared-cosine law, concentrating more steps near  $t = 0$ . This has been used in diffusion models for improved stability. With  $N \in \mathbb{N}$ :

$$t_i = \cos^2\left(\frac{\pi}{2} \cdot \frac{i}{N}\right), \quad i \in \{0, \dots, N\}, \quad t_N = 0. \quad (14)$$

**Chebyshev.** The Chebyshev schedule, derived from the nodes of Chebyshev polynomials of the first kind, minimizes polynomial interpolation error, and distributes steps more densely near boundaries:

$$t_i = \frac{1}{2} \left( \cos\left(\frac{\pi(i+0.5)}{N+1}\right) + 1 \right), \quad i \in \{0, \dots, N\}, \quad t_N = 0. \quad (15)$$

**EDM.** The EDM schedule [Karras et al., 2022] parameterizes the noise level  $\sigma$  using a power-law with exponent  $\rho \in \mathbb{R}^+$ . The interpolation is performed in  $\rho$ -space to allocate more steps where the generative process is most sensitive. For all experiments performed, we use the same parameters proposed by Karras et al. [2022], with  $\rho = 7$ ,  $\sigma_{\min} = 10^{-3}$ , and  $\sigma_{\max} = 1$ , and  $\sigma_{\min}, \sigma_{\max} \in \mathbb{R}^+$ :

$$\sigma_i = \left( \sigma_{\max}^{1/\rho} + \frac{i}{N} \left( \sigma_{\min}^{1/\rho} - \sigma_{\max}^{1/\rho} \right) \right)^{\rho}, \quad i \in \{0, \dots, N\}, \quad (16)$$

$$t_i = \frac{\sigma_i - \sigma_{\min}}{\sigma_{\max} - \sigma_{\min}}, \quad t_N = 0. \quad (17)$$

This formulation ensures a smooth transition between high noise and low noise.

### D.3 Dataset Details

For the datasets used, we follow the same training, validation, and test split used by [Tan et al., 2025a], where a single MCMC chain is decomposed into  $10^5$ ,  $2 \times 10^4$ , and  $10^4$  samples for training, validation, and test. The training and validation data are each taken as contiguous regions of the chain to simulate the realistic scenario where you have generated an MCMC trajectory and would like to use a Boltzmann Generator to continue generating samples. The data is split so that (after warmup) the first  $10^5$  samples are for train, the next  $2 \times 10^4$  are for validation, and the test samples are uniformly strided subsamples from the remaining MCMC chain. Earlier parts of the trajectory undersample certain modes enabling a biased training set that we attempt to debias through access to the energy function and SNIS. MD simulations were all run for 1  $\mu$ s, with a timestep of 1 fs, at temperatures of 300K, 310K, and 300K for alanine dipeptide, tri-alanine, and alanine tetrapeptide, respectively, in line with those conducted by Klein and Noe [2024]. The force fields used for all three molecules, in the same order, were the Amber ff99SBildn, Amber 14, and Amber ff99SBildn, in line with prior work by Tan et al. [2025a].

**Alanine Dipeptide.** The Ramachandran plots for the training and test split are provided in Fig. 5. Following Klein et al. [2023a], we up sample the lowest sampled mode (middle right) to make the problem easier for BGs. This also has the added benefit of testing what happens to FALCON when the training dataset is biased relative to the true distribution.

**Tri-alanine.** Similarly to alanine dipeptide, the Ramachandran plots for the training and test split are in Fig. 6. As there are two pairs of torsion angles that parameterize the system, there are two sets of Ramachandran plots included for each training and test. In tri-alanine, we can see that the training set actually misses a mode entirely ( $\psi_1 \approx \pi/3$ ), and undersamples this mode for  $\psi_2$  relative to the test set. This is a great test of finding a new mode (in  $\psi_1$ ) and correctly weighting a mode (in  $\psi_2$ ).

**Alanine Tetrapeptide.** For the tetrapeptide, there are three sets of Ramachandran plots each for training and test given the three pairs of torsions angles that parameterize the molecule, in Fig. 7.

**Hexa-alanine.** For hexa-alanine, we also include Ramachandran plots in Fig. 8. The first row shows training data, while the second row shows a held-out test set used to evaluate performance.

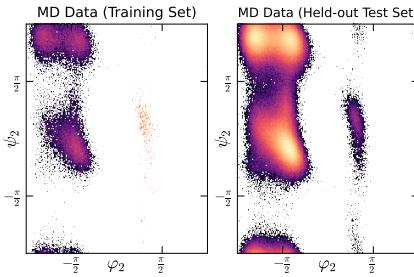


Figure 5: **Left:** Training data for alanine dipeptide; **Right:** Test data for alanine dipeptide.

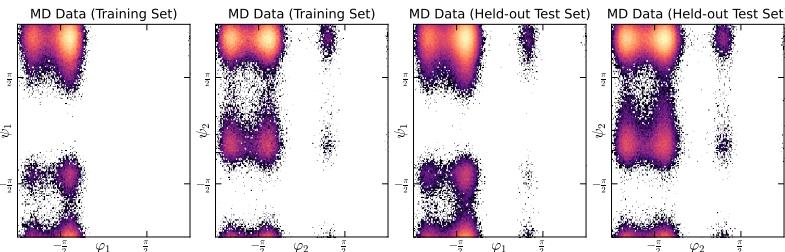


Figure 6: **Left and left center:** Training data for tri-alanine; **Right center and right:** Test data for tri-alanine.

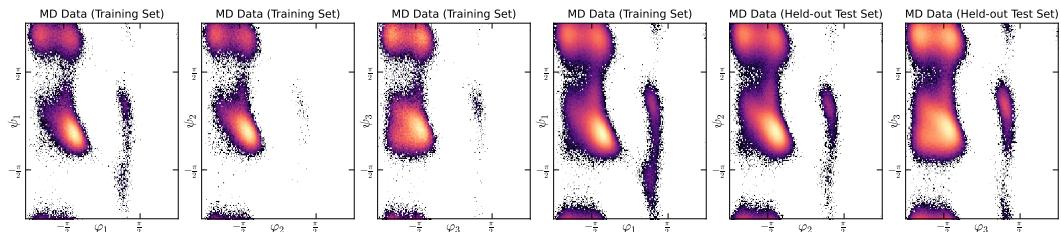


Figure 7: **First three:** Training data for alanine tetrapeptide; **Last three:** Test data for alanine tetrapeptide.

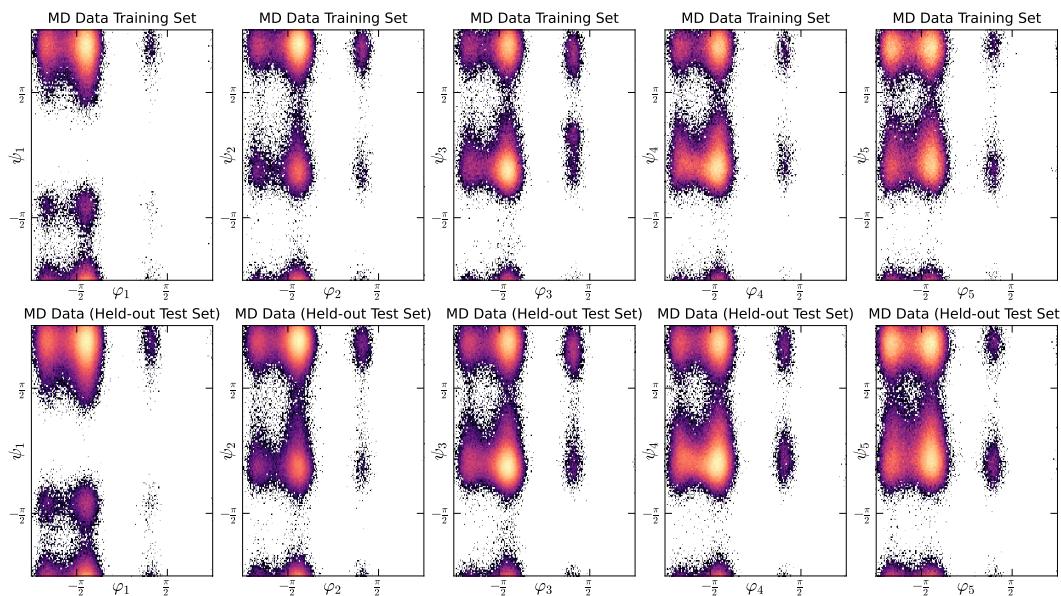


Figure 8: **First five:** Training data for hexa-alanine; **Last three:** Test data for hexa-alanine.

## E Metrics

**Effective Sample Size (ESS).** To quantify sampling efficiency, we compute the effective sample size (ESS) following Kish’s definition [Kish, 1957]. Given  $N \in \mathbb{N}$  generated particles with unnormalized importance weights  $\{w_i\}_{i=1}^N \subset \mathbb{R}^+$ , the ESS is normalized by  $N$  as:

$$\text{ESS}(\{w_i\}_{i=1}^N) \triangleq \frac{1}{N} \frac{1}{\sum_{i=1}^N w_i^2} \left( \sum_{i=1}^N w_i \right)^2. \quad (18)$$

The ESS reflects how many independent, equally-weighted samples would provide equivalent statistical power to the weighted sample. Higher ESS is desirable for more performant models.

**2-Wasserstein Energy Distance ( $\mathcal{E}\text{-}\mathcal{W}_2$ ).** To compare energy distributions, we compute the 2-Wasserstein distance between generated and reference samples from our MD dataset. For distributions  $p, q \in \mathcal{P}(\mathbb{R})$  over energy values, and  $\Pi(p, q)$  the set of admissible couplings, we define:

$$\mathcal{E}\text{-}\mathcal{W}_2(p, q)^2 \triangleq \min_{\pi \in \Pi(p, q)} \int_{\mathbb{R} \times \mathbb{R}} |x - y|^2 d\pi(x, y). \quad (19)$$

The  $\mathcal{E}\text{-}\mathcal{W}_2$  measures how closely the generated energy histogram aligns with that of the reference, with high sensitivity to structural accuracy due to bond-length-dependent energies. Since small perturbations in local structure induce large fluctuations in the energy distribution, this metric captures this variance, with lower values of  $\mathcal{E}\text{-}\mathcal{W}_2$  being more favourable for generative models.

**Torus 2-Wasserstein Distance ( $\mathbb{T}\text{-}\mathcal{W}_2$ ).** To assess structural similarity in torsional space, we compute a 2-Wasserstein distance over dihedral angles. For a molecule with  $L \in \mathbb{N}$  residues, define the dihedral vector as:

$$\text{Dihedrals}(x) = (\phi_1, \psi_1, \dots, \phi_{L-1}, \psi_{L-1}) \in [0, 2\pi)^{2(L-1)}. \quad (20)$$

The cost on the torus accounts for periodicity of angles:

$$c_{\mathcal{T}}(x, y)^2 = \sum_{i=1}^{2(L-1)} \left[ (\text{Dihedrals}(x)_i - \text{Dihedrals}(y)_i + \pi) \bmod 2\pi - \pi \right]^2. \quad (21)$$

The torus 2-Wasserstein distance between two distributions  $p, q \in \mathcal{P}([0, 2\pi)^{2(L-1)})$  is then:

$$\mathbb{T}\text{-}\mathcal{W}_2(p, q)^2 \triangleq \min_{\pi \in \Pi(p, q)} \int c_{\mathcal{T}}(x, y)^2 d\pi(x, y). \quad (22)$$

This captures macroscopic conformational differences while respecting angular periodicity. The  $\mathbb{T}\text{-}\mathcal{W}_2$  provides a more global assessment of performance, for instance revealing missed conformational modes that would not be captured by the  $\mathcal{E}\text{-}\mathcal{W}_2$ .

## F FALCON Algorithm

We detail the FALCON training algorithm below.

---

### Algorithm 1: Training FALCON

---

**Input:** Sampleable  $p_0$  and  $p_1$ , regularization weight  $\lambda_r$ , network  $u_\theta$

**Output:** The trained network  $u_\theta$

**while** training **do**

$$\begin{aligned} & (x_0, x_1) \sim p_0 \times p_1; \\ & x_s \leftarrow sx_1 + (1-s)x_0; \\ & v_s \leftarrow x_1 - x_0; \\ & u_\theta, \frac{\partial u_\theta}{\partial s} \leftarrow \text{jvp}(u_\theta, (x_s, s, t), (v_s, 1, 0)); \\ & u_{\text{tgt}} \leftarrow v_s - (t-s)\frac{\partial u_\theta}{\partial s}; \\ & \hat{x}_t \leftarrow x_s + (t-s)u_\theta; \\ & \hat{x}_s \leftarrow \hat{x}_t + (s-t)u_\theta(\hat{x}_t, t, s); \\ & \mathcal{L}(\theta) \leftarrow \|u - \text{stopgrad}(u_{\text{tgt}})\|^2 + \lambda_r \|x_s - \hat{x}_s\|^2; \\ & \theta \leftarrow \text{update}(\theta, \nabla_\theta \mathcal{L}(\theta)); \end{aligned}$$

**return**  $u_\theta$ ;

---

## G Additional Experiments

**Impact of Inference Schedules.** In the few step regime, performance can be significantly impacted by the choice of inference schedule. We run ablations on various schedules for alanine dipeptide with 8-steps, summarizing the results in Fig. 9. We note that the EDM scheduler substantially outperforms all other schedulers, in agreement with observations from the diffusion literature [Karras et al., 2022]; sampling more points near the data distribution proves beneficial in aiding generative performance as the variance of the flow field is higher closer to the target distribution. For all reported results, we use the EDM scheduler. In Appendix D.2, we provide additional details regarding scheduler definitions, inference setups, and parameter selection for EDM.

### G.1 Details for Figure Generation

In Fig. 1, we compare performance on the  $\mathbb{T}\mathcal{W}_2$  metric for alanine dipeptide between FALCON and our continuous-time DiT CNF. For FALCON, accuracy is controlled by the number of inference steps (1–8). For the DiT CNF, we consider three settings: (1) adaptive step Dormand–Prince 4(5) with exact Jacobian trace evaluation, varying atol/rtol from  $10^{-1}$  to  $10^{-5}$ ; (2) the same tolerances but instead using the Hutchinson trace estimator, trading performance for faster likelihoods with higher variance; and (3) fixed-step Euler integration with step sizes from 4 to 256 with exact Jacobian traces, where the upper bound is chosen to roughly match the number of function evaluations needed for Dopri5. The models and training configurations used are presented in Appendix D.1.

In Fig. 2, we demonstrate the energy distributions for unweighted and re-weighted samples for our most performant FALCON Flows. For alanine dipeptide, tri-alanine, alanine tetrapeptide, and hexa-alanine, 4, 8, 8, and 16 steps were used for figure generation. Energy distributions reveal microscopic detail, as marginal changes in local atomic position can have significant impacts on total energy. The best models for each molecular system were used, with pertinent details on model size, training configurations, and hyperparameters detailed in Appendix D.1. Similarly, in ??, we show the proposal and re-weighted sample estimates for alanine dipeptide, and demonstrate how an increasing number of steps, improves the energy distribution.

In Fig. 3, we compare the performance between SBG’s discrete NFs and FALCON, illustrating that despite more samples increasing performance across metrics, the approach is still unable to reach FALCON’s performance. For SBG, we specifically take their best model weights for the TarFlow architecture from: <https://github.com/transferable-samplers/transferable-samplers>, and draw  $N_s \in \{10^4, 2 \times 10^4, 5 \times 10^4, 10^5, 2 \times 10^5, 5 \times 10^5, 10^6, 2 \times 10^6, 5 \times 10^6\}$  samples three different times (the error bars are representative of the three draws) and evaluate  $\mathcal{E}\mathcal{W}_2$  for each set.

In Fig. 4, we investigate how the strength of regularization impacts performance on ESS and  $\mathcal{E}\mathcal{W}_2$ . Specifically, we demonstrate that small amounts of regularization enable generative performance but impede invertibility, while too much regularization detrimentally impacts sample quality. This trade-off leads to an optimum on both metrics. We investigate  $\lambda_r \in \{10^0, 10^1, 10^2, 10^3, 10^4, 10^5\}$ , and conclude upon  $\lambda_r = 10^1$ . To ascertain the optimal  $\lambda_r$ , we ran these experiments on alanine dipeptide, using the same model details and configurations highlighted in Appendix D.1.

Fig. 9 demonstrates the sensitivity the inference schedule plays on generative performance. For this figure, we took our best models—trained with the details presented in Appendix D.1—and ran inference using all three trained seeds to generate uncertainties for each inference schedule. We used our 8-step system, as inference scheduler choice is less important the fewer steps are used.

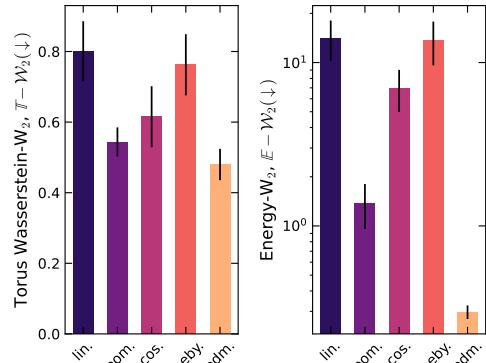


Figure 9: Performance vs. choice of inference schedule. The left plot shows the Torus Wasserstein-W<sub>2</sub> metric, and the right plot shows the Energy-W<sub>2</sub> metric. Both plots compare five inference schedules: lin., geom., cos., cheby., and edm. The edm schedule consistently outperforms the others in both metrics.

## G.2 Proof of Invertibility

In our loss, we use a cycle-consistency term that regularizes training to promote numerical invertibility. We see in Fig. 4, the introduction of this modified loss aids generative performance; however, the forward-backward reconstruction yields errors on the order of  $10^{-2}$ , indicating an approximately, but not entirely invertible model. To prove that the flow is invertible (but the inverse is challenging to discover during training), we train an auxiliary FALCON Flow in the reverse direction. We use the same network to parameterize the reverse flow as described in Appendix D.1, with the same training configuration and hyperparameter set. Next, we freeze the forward model (which goes from latents  $\rightarrow$  data), generate synthetic data ( $2 \times 10^5$  prior-target sample pairs) and train the auxiliary model on these reflow targets to learn the mapping from data back to latents. In Fig. 10, we illustrate the loss curve of the trained auxiliary FALCON. To evaluate invertibility, we draw i.i.d. samples from our prior distribution, pass them through the frozen forward model, and test the auxiliary reverse model on these unseen latents. We evaluate an  $\ell_2$  error on the recovered latents to find reconstruction within  $10^{-4}$  after 1000 epochs matching the reconstruction accuracy of discrete NFs that are invertible by design. These results support the claim that the learned flow is indeed invertible.

## G.3 Performance Against Number of Function Evaluations

Fig. 11 highlights the efficiency of FALCON in terms of number of function evaluations. We see that FALCON achieves the same torus 2-Wasserstein performance as our DiT CNF with Dopri5, while requiring over two orders of magnitude fewer function evaluations. Whereas Fig. 1 quantified efficiency in terms of inference time, here we measure the number of function evaluations against both fixed-step solvers (Euler with 4–256 steps) and adaptive solvers (Dopri5 run until reaching target tolerances  $atol = rtol \in \{10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}\}$  for Hutchinson and  $\in \{10^{-2}, 10^{-3}, 10^{-4}, 10^{-5}, 10^{-6}\}$  for exact).

Although Hutchinson’s trace estimator yields a speedup relative to exact Jacobian computation for CNFs, as seen in Fig. 1, the number of function evaluations needed is higher than the exact Jacobian computation. In either setting, however, FALCON still remains substantially faster. Even with a 4-step solver, FALCON matches the accuracy of DiT CNFs using Hutchinson at  $atol = rtol = 10^{-5}$ , with an 8-step solver nearly matching the performance of a DiT CNF with Dopri5 set to an  $atol = rtol = 10^{-6}$ .

## G.4 Ramachandran Plots

**Alanine Dipeptide.** We demonstrate FALCON’s capacity to learning global features through the Ramachandran plots in Fig. 12 for alanine dipeptide. We include both the held-out test set and the learned model’s map, showcasing its ability to debias the data and capture the true MD distribution. Specifically, the undersampled  $\phi$  mode in the training data is correctly upweighted in the learned model’s predictions, indicating accurate likelihood estimates from the learned flow.

**Tri-alanine.** Similarly, we show the Ramachandran plots for tri-alanine in Fig. 13 exhibiting similar behaviour. Most conformations are correctly captured, with some modes being underweighted.

**Alanine Tetrapeptide.** Next, we show the Ramachandran plots for alanine tetrapeptide in Fig. 14.

**Hexa-alanine.** Finally, we show the Ramachandran plots for hexa-alanine in Fig. 15.

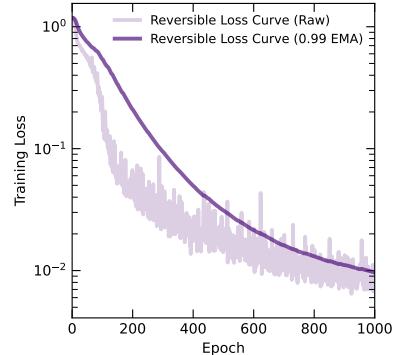


Figure 10: Auxiliary model loss from reflow target training on forward flow.

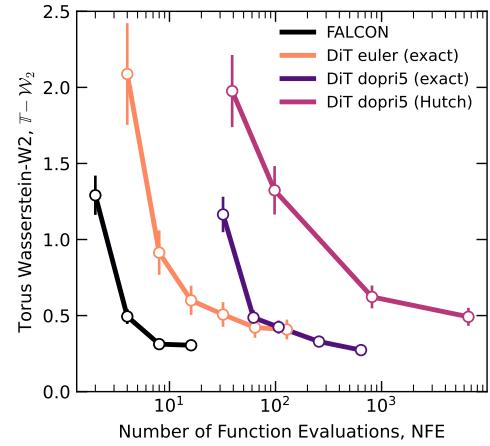


Figure 11: Performance of FALCON vs. our DiT CNFs as a function of NFEs.

Even with a 4-step solver, FALCON matches the accuracy of DiT CNFs using Hutchinson at  $atol = rtol = 10^{-5}$ ,

with an 8-step solver nearly matching the performance of a DiT CNF with Dopri5 set to an  $atol = rtol = 10^{-6}$ .

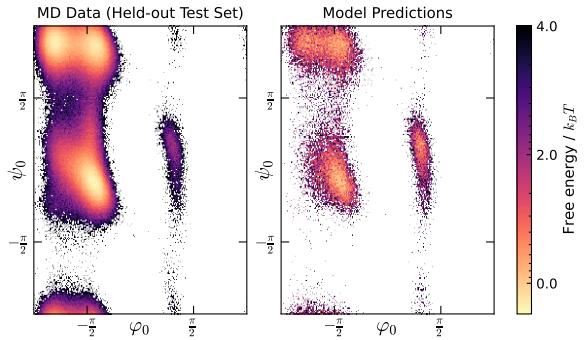


Figure 12: **Left:** Test data for alanine dipeptide; **Right:** FALCON’s angular predictions for alanine dipeptide.

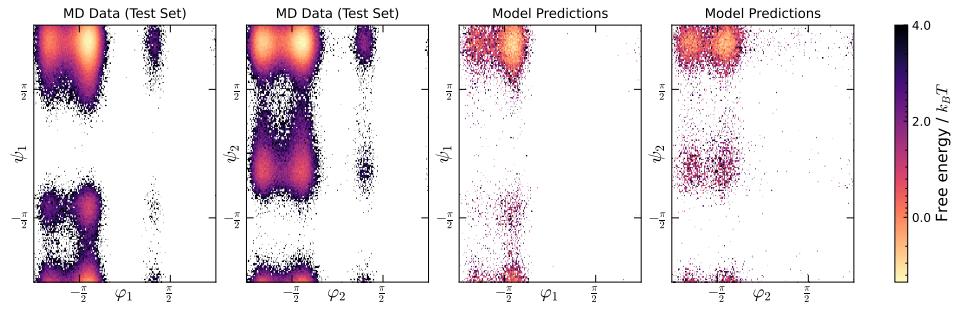


Figure 13: **Left and left center:** Test data for tri-alanine; **Right and right center:** FALCON’s angular predictions for tri-alanine.

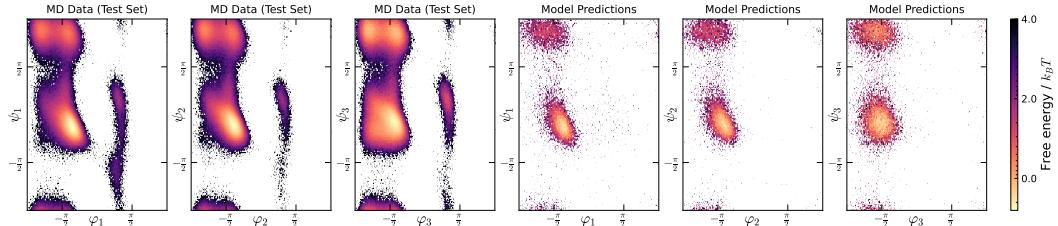


Figure 14: **First three:** Test data for alanine tetrapeptide; **Last three:** FALCON’s angular predictions for alanine tetrapeptide.

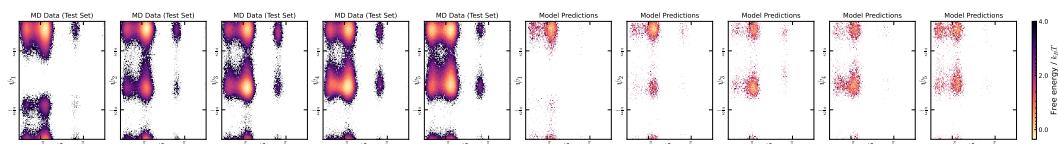


Figure 15: **First five:** Test data for hexa-alanine; **Last five:** FALCON’s angular predictions for hexa-alanine.